

Redback Racing – Weather Station Cloud Integration (2025 T3)

1. Introduction

Currently, Redback Racing captures and streams on-car telemetry through AWS services. To provide engineers with a better context, weather data must be integrated into the cloud platform. A trackside weather station will measure temperature, humidity, wind, and track surface conditions. The goal is to design a solution that is **secure, scalable, low-latency, and cost-effective**, while fitting cleanly into the existing AWS architecture.

2. Proposed Architecture

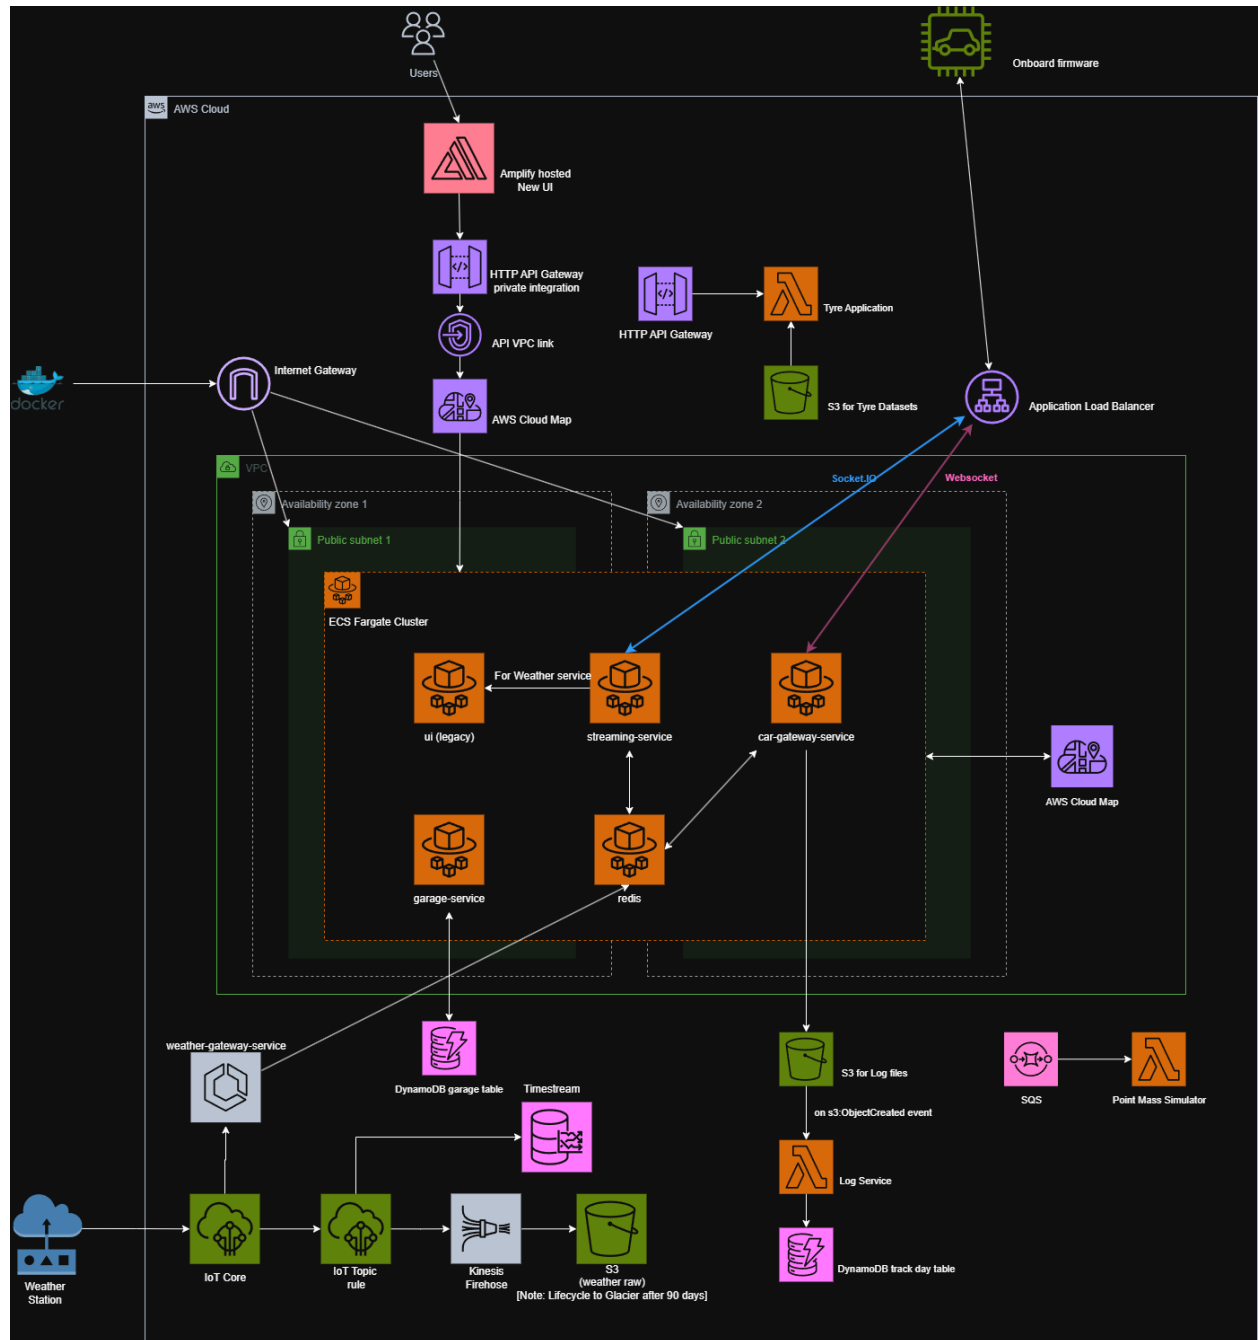
The proposed flow is:

Weather Station → AWS IoT Core (MQTT/TLS) → IoT Rule → Kinesis Firehose → S3 (raw) + Timestream (queries) → ECS Fargate weather-gateway-service → Redis → streaming-service → UI (1)

Explanation of each device:

- Weather Station: publishes JSON sensor data over MQTT (topic: weather/{track_id}/{station_id}/metrics). (2)
- IoT Core: authenticates devices with X.509 certificates, terminates TLS (transport layer security), and routes messages.
- IoT Rules + Firehose: fan-out telemetry to both S3 for raw archival and Timestream for structured time-series queries.
- ECS weather-gateway-service: subscribes to MQTT topics, forwards sanitised data into Redis, and integrates with the existing streaming-service and UI.

The diagram below shows the addition of the architecture:



3. Protocol Choice

The weather station uses MQTT over TLS 1.2+. MQTT is a lightweight publish/subscribe protocol optimised for IoT, offering delivery guarantees through Quality of Service (QoS) levels (3). The QoS level that we will choose is QoS 1, which will ensure a reliable message delivery even on unstable LTE/Wi-Fi links, with duplicates handled downstream. This is more efficient than HTTP polling and aligns with AWS IoT Core's native features (4).

4. Security

- Device identity: each station has an X.509 certificate registered in IoT Core with least-privilege policies (e.g., can only publish to `weather/{track}/{station}/{#}`).
 - Encryption: TLS in transit; S3 SSE-KMS and Timestream encryption at rest.
 - Networking: ECS tasks run in private VPC subnets with NAT; no direct internet access.
 - Audit: CloudWatch for metrics/logs; CloudTrail for API activity.
 - Secrets management: IAM task roles for ECS, no static credentials embedded in devices.
-

5. Scalability & Reliability

- Managed ingest: IoT Core and Firehose scale elastically to thousands of messages/sec.
 - Firehose buffering: batches efficiently before writing to S3.
 - Storage tiers: S3 lifecycle policies archive data to Glacier after 90 days.
 - Retention: Timestream memory store holds hot data (72h), magnetic store retains 1 year.
 - Reliability: devices use local ring buffers to store readings during outages, retry on reconnect, and publish with QoS 1. Firehose has an error S3 prefix for failed records.
-

6. Performance

- Live path: MQTT → weather-gateway-service → Redis → streaming-service → UI. This supports sub-second visibility in dashboards.

- Batch path: Firehose → S3 + Timestream for analytics. Optional Firehose transformations can convert to Parquet for Athena queries.
 - Caching: Redis stores latest readings (weather:latest) with TTL, ensuring fast O(1) lookups.
-

7. Cost Analysis

- IoT Core: charged per million messages; low cost at weather station scale.
 - Kinesis Firehose: pay per GB ingested and transformed; minimal overhead.
 - S3: cost-efficient long-term storage; lifecycle to Glacier further reduces cost.
 - Timestream: billed by ingest and query; cheaper than self-hosting a time-series DB.
 - ECS Fargate: billed per vCPU and GB-hour; only pays when containers are running.
This design avoids higher-cost services like MSK/Kafka, which are unnecessary at the current scale.
-

8. Integration with Existing System

The car telemetry path remains unchanged. The weather path adds one ECS service (weather-gateway-service) into the existing Fargate cluster, forwarding live data into Redis. This allows the streaming service and UI to display weather alongside car telemetry with minimal changes. Historical data in S3/Timestream can be joined with car logs for analysis.

9. Terraform (Infrastructure as Code) Examples

```
provider "aws" { region = "ap-southeast-2" }
```

```
resource "aws_s3_bucket" "weather_raw" {  
  bucket = "rbr-weather-raw-${var.env}"  
}
```

```
resource "aws_s3_bucket_lifecycle_configuration" "weather_raw_lc" {  
  bucket = aws_s3_bucket.weather_raw.id  
  rule {
```

```

id    = "tiering"
status = "Enabled"
transition { days = 30 storage_class = "INTELLIGENT_TIERING" }
transition { days = 90 storage_class = "GLACIER" }
}
}

resource "aws_timestreamwrite_database" "weather" {
  database_name = "rbr_weather_${var.env}"
}

resource "aws_timestreamwrite_table" "measurements" {
  database_name = aws_timestreamwrite_database.weather.database_name
  table_name    = "measurements"
  retention_properties {
    memory_store_retention_period_in_hours = 72
    magnetic_store_retention_period_in_days = 365
  }
}

```

10. Risks & Mitigations

- Connectivity loss → station caches locally, retries with QoS 1.
 - Clock drift → devices sync with NTP; server-side timestamps as backup.
 - Schema evolution → versioned MQTT topics (v1/metrics), JSON schema validation in gateway.
 - Scaling edge cases → S3 lifecycle and Timestream retention prevent runaway costs.
-

11. Conclusion

This design securely integrates trackside weather telemetry into Redback's AWS system. It leverages managed services (IoT Core, Firehose, S3, Timestream, ECS Fargate) to minimise operational overhead, ensure low-latency live updates for engineers, and provide a durable, queryable history for analysis. The approach balances security, scalability, performance, and cost, while fitting neatly into the existing ECS + Redis + UI ecosystem.

12. Appendix.

(1) Some other alternatives that were considered but not picked:

1. API Gateway (WebSocket/HTTP) → ECS service → S3/DB (no IoT Core)
 - a. Pros: Fewer AWS services; familiar REST/WebSocket model.
 - b. Cons: No device identity/cert provisioning workflow; heavier protocol for embedded; you reinvent features IoT Core already solves (authN, topic routing, QoS).
 - c. Why not: We need secure device onboarding + low-power telemetry. MQTT/IoT Core is purpose-built.
2. IoT Core → Lambda → S3/Timestream (no Firehose, no ECS gateway)
 - a. Pros: Really simple serverless fan-out; fewer services; easy transforms in Lambda.
 - b. Cons: Harder to feed live data into the existing Redis/UI without introducing a push channel; risk of Lambda cold starts for near-real-time UX.
 - c. Why not: We want a guaranteed low-latency live path into Redis while still capturing a durable history.
3. IoT Core → Kinesis Data Streams → consumers (ECS/Lambda) → S3/Timestream
 - a. Pros: Strong real-time streaming semantics, multiple consumer apps, replays.
 - b. Cons: More Ops and cost than Firehose for a single weather feed; extra complexity (shards, scaling policies).
 - c. Why not: Weather volume is modest; Firehose is simpler/cheaper for ingestion to storage.
4. Use DynamoDB (time-series table) instead of Timestream
 - a. Pros: Familiar, versatile, predictable cost.
 - b. Cons: You must design/tune TS schema, TTL, and queries; lacks native TS functions/retention tiers.
 - c. Why not: Timestream is built for time-series (ingest/retention/queries) and pairs well with S3 for raw.
5. Run our own TSDB (InfluxDB/Timescale) on ECS/EKS

- a. Pros: Full control, rich TS features.
 - b. Cons: You own backups, scaling, HA, patching → higher operational burden.
 - c. Why not: Managed services reduce toil; scope is an assessment with limited time.
6. MSK/Kafka
- a. Pros: Great for large multi-topic, multi-consumer streaming ecosystems.
 - b. Cons: Expensive and operationally heavy for one station's telemetry.
 - c. Why not: Overkill here.
7. IoT Greengrass or local edge compute
- a. Pros: Local rules/aggregation if connectivity is poor; can pre-filter data.
 - b. Cons: Extra device complexity; not needed unless bandwidth/offline operation is a hard requirement.
 - c. Why not (for now): Keep edge simple; add later if requirements expand.
8. Skip weather-gateway-service and push UI directly from IoT Core (MQTT over WebSocket)
- a. Pros: Fewer components between device and UI.
 - b. Cons: Auth/session management in browser, topic authorization complexity, bypasses existing Redis/streaming-service pattern.
 - c. Why not: We want to reuse the same live data path (Redis + streaming-service) the team already operates.

Summary on why I picked what I picked:

- We chose IoT Core + Firehose + S3/Timestream for a secure, managed, low-ops backbone, and a small ECS gateway to slot the feed into the existing Redis/UI for real-time UX.
- The alternatives either add ops/cost (Kafka, DIY TSDB), weaken device identity/efficiency (pure API GW/HTTP), or lose low-latency UI fit (serverless-only fan-out).

(2) Example JSON code:

```
{  
  "ts": 1725345600.123,  
  "air_temp_c": 28.6,
```

```
"humidity_pct": 54.2,  
"wind_speed_ms": 7.3,  
"wind_dir_deg": 210,  
"track_temp_c": 35.9,  
"station_id": "ws-01",  
"track_id": "smp"  
}
```

(3) QoS → Quality of Service level:

- QoS 0 → “At most once”
 - The sender sends the message once, no acknowledgment.
 - If the network drops, the message is lost.
 - Lowest overhead, fastest.
 - Use when occasional loss is fine (e.g. live temperature updates every second).
- QoS 1 → “At least once”
 - Sender retries until it gets an acknowledgment from the receiver.
 - The message might arrive more than once (duplicates possible).
 - Most common in IoT — reliable but still lightweight.
 - Good for telemetry where you’d rather have duplicates than lose data.
- QoS 2 → “Exactly once”
 - Sender and receiver go through a 4-step handshake.
 - Guarantees delivery without duplicates.
 - Highest overhead and slowest.
 - Use only for critical transactions (e.g. turning a relay on/off).

(4) Reason why it is more efficient:

MQTT was designed for IoT and telemetry, whereas HTTP was designed for request–response applications. MQTT is more efficient than HTTP polling for several reasons:

1. Connection model

- HTTP: Every request/response opens a new TCP/TLS connection unless you manage keep-alive.

- MQTT: Maintains a long-lived TCP/TLS connection with very low overhead.

2. Message size

- HTTP: Includes verbose headers (hundreds of bytes per request).
- MQTT: Minimal fixed header (2 bytes), so small sensor payloads aren't drowned out by metadata.

3. Communication style

- HTTP: Client must poll ("ask repeatedly") for new data, wasting bandwidth if nothing has changed.
- MQTT: Uses publish/subscribe, so devices only send when data changes, and subscribers get updates instantly.

4. Delivery guarantees

- HTTP: Either you get the response or not, but no fine-grained control.
- MQTT: Built-in QoS levels (0, 1, 2) let you balance reliability vs. performance.

5. Power consumption

- For IoT devices on battery, HTTP polling burns power keeping radios active.
- MQTT's lightweight, event-driven nature is more energy-efficient.

Summary:

MQTT over TLS reduces bandwidth, latency, and power consumption compared to HTTP polling, making it far better for IoT telemetry like weather stations.